

Lecture 3

Neural nets and backpropagation

1. NER - Named Entity Recognition

- task of labeling words in text by **word tokens** for ex. name, date etc
- it is done by: taking each word in its fixed **context window** forming word vectors, then running through **classifier** which gives out the probability of its label
- **The classifier:**
 - so if the context window size is 5, then the input would be a 5D-vector \mathbf{x}
 - and it is passed through a neural network
 - x (input) $\in \mathbb{R}^{5d}$ — input vector
 - $h = f(Wx + b)$ — hidden layer (f is activation function)
 - $s = u^T h$ — score
 - $J_t(\theta) = \sigma(s) = \frac{1}{1+e^{-s}}$ — predicted probability (sigmoid) close to 1 if high score
- **Gradients of matrix:**
 - **Jacobian Matrix: Generalization of the Gradient**

- Given a function with **m outputs** and **n inputs**
$$\mathbf{f}(\mathbf{x}) = [f_1(x_1, x_2, \dots, x_n), \dots, f_m(x_1, x_2, \dots, x_n)]$$
- It's Jacobian is an **$m \times n$ matrix** of partial derivatives

$$\frac{\partial \mathbf{f}}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

$$\left(\frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right)_{ij} = \frac{\partial f_i}{\partial x_j}$$

- 回味多元微积分...
- **Elementwise function:** has some special properties in which $h_i = f(x_i)$ and $\frac{\partial h}{\partial x_i}$ is computed
 - its **Jacobian matrix** is a diagonal matrix
- **simple properties**
 - $\frac{\partial}{\partial x}(Wx + b) = W$
 - $\frac{\partial}{\partial b}(Wx + b) = I$
 - $\frac{\partial}{\partial u}(u^T h) = h^T$ just if you treat a vector as row. so for word vecs, you treat it normally as column vector

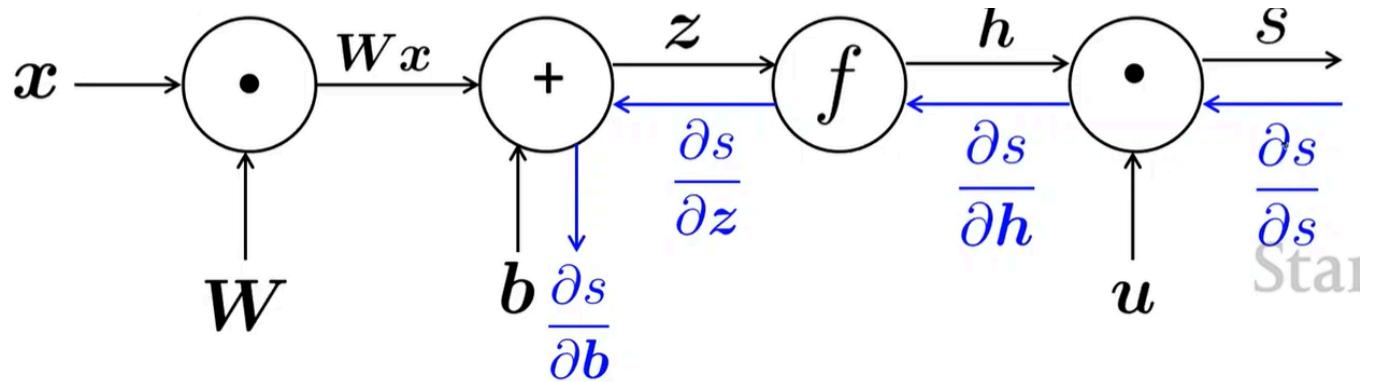
2. Neural Nets

- **Working out the gradient of score function:**
 - $s = u^T h$
 - $h = f(z)$
 - $z = Wx + b$
 - so to **update b** , we compute the partial derivative of s with respect to b
 - $\frac{\partial s}{\partial b} = \frac{\partial s}{\partial h} \frac{\partial h}{\partial z} \frac{\partial z}{\partial b}$
 - Where $z = Wx + b$, applying chain rule:

- $= u^T \text{diag}(f'(z))I$
- $= u^T \circ f'(z)$ where the \circ is hadamard product, makes it into a $1 \times n$ **row vector**
- because b must be as the same shape as Wx
- **The method to update W is similar**
- Similarly, the first 2 partial derivatives are **identical**, so we can remove the need of recomputing. **Taking it as δ**
- **Shape Convention:**
 - because for deriving in terms of W , $\frac{\partial s}{\partial W}$, the score needs to be a scalar number, so there is $n \times m$ inputs, and 1 output so the **jacobian is a $1 \times nm$ row vector**.
 - but in **shape convention** you convert the $1 \times nm$ into the shape of **W**
 - and eventually you use the same $n \times m$ matrix to update the new W
- **Actual and rigorous derivation:**
 - so $\frac{\partial s}{\partial W} = \delta \frac{\partial z}{\partial W}$
 - dimension of δ is $n \times 1$, because W is $n \times m$, x as input, is having $m \times 1$.
 - and $\frac{\partial z}{\partial W} = x$
 - proof of this derivation
 - W_{ij} maps the input x_j to the neuron i
 - so we have $z_i = \sum_{k=1}^m W_{ik}x_k + b_i$
 - so in each derivation, only one term remains
 - so to fit with the dimension $\frac{\partial s}{\partial W} = \delta x^T$ (typo in the lecture???)
- Overall, the gradient has to be the same shape as the parameter

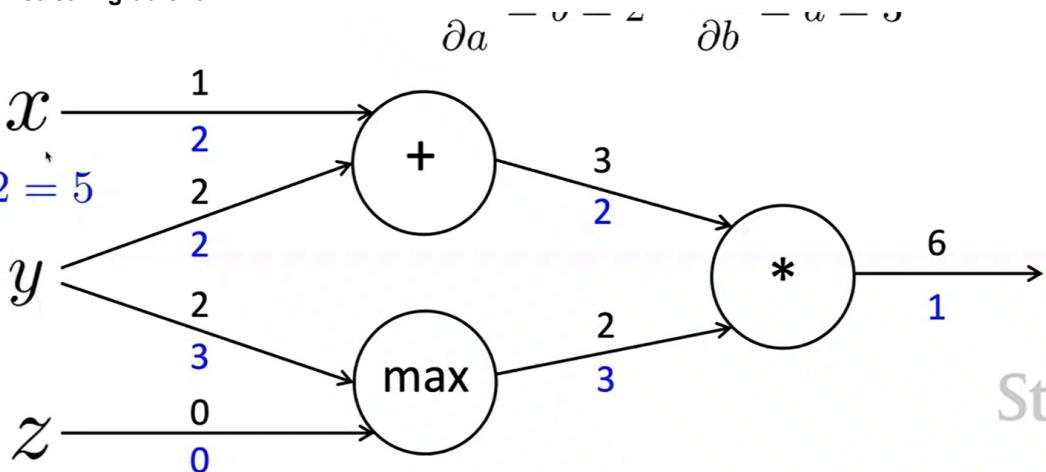
3. Backpropagation

- process of sending back gradients, to update the parameters of the models to learn, to give more accurate outcome and reduce loss.



- **Additionally**, we can make use of the **chain rule**, where with an **upstream gradient**, and calculating the **local gradient** we could easily find out the **downstream gradient**.

$$\begin{aligned}\frac{\partial f}{\partial x} &= 2 \\ \frac{\partial f}{\partial y} &= 3 + 2 = 5 \\ \frac{\partial f}{\partial z} &= 0\end{aligned}$$



- and eventually for here, the gradient on the input represents how much changes you make on the input, that gradient times the change would be the change of final output

Compute gradient wrt each node using gradient wrt successors

$\{y_1, y_2, \dots, y_n\}$ = successors of x

$$\frac{\partial z}{\partial x} = \sum_{i=1}^n \frac{\partial z}{\partial y_i} \frac{\partial y_i}{\partial x}$$

- **Intuitions**

- '+' distributes the upstream gradient(the step after) to both summands (no need of local gradient)
- 'max' routes the upstream gradient (chooses one)
- '*' switches the upstream gradient for cases of xy (uses the value of other route, after derivating in terms of one)

- **Autoamatic Differentiation:** calculating the corresponding gradients before hand.