


Assignment 1

我做的是CS224N 2025秋季

总体来说第一个assignment还是挺容易的，就是配置浪费了好多时间，要注意的就是：


- 用Anaconda配置默认的python版本是3.13, 就是下载时也没有任何让你修改版本的选项，然后这导致的就是import 的 gensim数据库无法导入，因为python 3.13不兼容，需降级到3.12.
- 我去谷歌到处搜别人做过的作业，发现好像就2025 Manning 教授才开始使用gensim库。这里降级的操作可以参考这位博主的方法：<https://blog.csdn.net/szjia/article/details/147540541>
- 虽然整体都不难，但是第一个assignment所让我接触到对于我来说很新颖的东西。这个assignment的主旨就是利用庞大文字段落创建word vector
- wordvector：就是要把每一个单词作为主单词依次枚举，然后看周围一定半径内出现的其他单词次数。这就可以造出一个 $n \times n$ 矩阵
- 随后再用TruncatedSVD库来调用线性代数中SVD算法来对这个矩阵进行压缩
- 最后使用matplotlib, 来根据每个wordvec 的坐标在一个二维坐标系下标注他们位置，由于之前算过每个单词周围出现的单词的概率，所以不同的单词周围出现的其他单词若相似度很高，这样，用法、即意义相似的单词就会聚集在一起
- 下面即是作业中算出word vector 后用matplotlib 点出的graph



Lecture 1

Beginning of the lecture

- Word2Vecs: are a way of representing words as vectors
 - it contains of a **center word** and a **context window** which are the surrounding words of the center word, thus giving a plot for a n-dimensional diagram, enabling us to see similar words being grouped together
 - Word vectors, also referred as word embeddings 演示中为4到5维的，实际上使用300多维的向量



Likelihood functions

For each position $t = 1, \dots, T$, predict context words within a window of fixed size m , given center word w_t . Data likelihood: 大概的意思

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m \\ j \neq 0}} P(w_{t+j} | w_t; \theta)$$

theta contains all the wordvecs - word embeddings
 θ is all variables to be optimized

sometimes called a *cost* or *loss* function

$\log(i=1 \prod nai) = i=1 \sum n \log(ai)$

The objective function $J(\theta)$ is the (average) negative log likelihood:

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

Minimizing objective function \Leftrightarrow Maximizing predictive accuracy

- The first Likelihood function coming from that j being index for the context words, so the first product sum representing the probability for a single center word with all the context word
- Then, t , representing the center word, so the other product sum representing for looping all the center words, and the size of the context window being from $-m$ to m
- Additionally, we would like to maximize the likelihood function, which is the same as **minimizing the negative log likelihood function - the cost function $J(\theta)$**
- And we would like to maximize the likelihood because the model needs to consistently give **high probability to the real context words** that appear around each center word.
-

$$P(o|c) = \frac{\exp(u_o^T v_c)}{\sum_{w \in V} \exp(u_w^T v_c)}$$

- and the P function itself, is a **softmax function**
- One word, has **2 corresponding word vecs**, one for the center word, and one for the context word
- $u_o^T v_c$ dot product resembling similarity, and exponentiation (positive) and normalizing to image it to 0 to 1 to make it a probability.

- the $J(\theta)$ also refers to the **cost/loss function**, because it shows how bad the model's prediction is compared to what we want
- however, is inefficient due to the need of calculating all the expectancies (dot products) for every single word each and every time for the denominator

But the dimensions are too large?!

the range is $2 * d * V!!$, so there is a need of gradient

- so taking the partial derivative of the $J(\theta)$ function in terms of v_c so optimizing it can predict its context words better
- ..简单的偏导计算，略
- ends up in = **observed - expected** which is exactly what we need to minimize, and the gradient refers to the direction of the steepest ascent